

# image-classification

September 20, 2023

```
[ ]: model_checkpoint = "microsoft/swin-tiny-patch4-window7-224" # pre-trained model
      batch_size = 32 # batch size for training and evaluation
```

```
[1]: !pip install -q datasets transformers
```

```
519.6/519.6
kB 5.4 MB/s eta 0:00:00
7.6/7.6 MB
26.2 MB/s eta 0:00:00
115.3/115.3 kB
15.0 MB/s eta 0:00:00
194.1/194.1 kB
22.5 MB/s eta 0:00:00
134.8/134.8 kB
16.3 MB/s eta 0:00:00
294.9/294.9 kB
32.3 MB/s eta 0:00:00
7.8/7.8 MB
68.9 MB/s eta 0:00:00
1.3/1.3 MB
76.3 MB/s eta 0:00:00
```

```
[2]: from huggingface_hub import notebook_login
```

```
notebook_login()
```

```
VBox(children=(HTML(value='<center> <img\nsrc=https://huggingface.co/front/
assets/huggingface_logo-noborder.svg...</center>'))
```

```
[ ]: %%capture
      !sudo apt -qq install git-lfs
      !git config --global credential.helper store
```

```
[ ]: from transformers.utils import send_example_telemetry

      send_example_telemetry("image_classification_notebook", framework="pytorch")
```

### 0.0.1 Loading the dataset

```
[ ]: from datasets import load_dataset

# load a custom dataset from local/remote files or folders using the
# ↳ ImageFolder feature

# option 1: local/remote files (supporting the following formats: tar, gzip,
# ↳ zip, xz, rar, zstd)
dataset = load_dataset("imagefolder", data_files="https://madm.dfki.de/files/
# ↳ sentinel/EuroSAT.zip")

# note that you can also provide several splits:
# dataset = load_dataset("imagefolder", data_files={"train": ["path/to/file1",
# ↳ "path/to/file2"], "test": ["path/to/file3", "path/to/file4"]})

# note that you can push your dataset to the hub very easily (and reload
# ↳ afterwards using load_dataset)!
# dataset.push_to_hub("nielsr/eurosat")
# dataset.push_to_hub("nielsr/eurosat", private=True)

# option 2: local folder
# dataset = load_dataset("imagefolder", data_dir="path_to_folder")

# option 3: just load any existing dataset from the hub, like CIFAR-10,
# ↳ FashionMNIST ...
# dataset = load_dataset("cifar10")
```

Using custom data configuration default-0537267e6f812d56

Downloading and preparing dataset image\_folder/default to /root/.cache/huggingface/datasets/image\_folder/default-0537267e6f812d56/0.0.0/ee92df8e96c6907f3c851a987be3fd03d4b93b247e727b69a8e23ac94392a091...

Downloading data files: 0it [00:00, ?it/s]

Downloading data files: 0%| | 0/1 [00:00<?, ?it/s]

Downloading data: 0%| | 0.00/94.3M [00:00<?, ?B/s]

Extracting data files: 0%| | 0/1 [00:00<?, ?it/s]

Generating train split: 0 examples [00:00, ? examples/s]

Dataset image\_folder downloaded and prepared to /root/.cache/huggingface/datasets/image\_folder/default-0537267e6f812d56/0.0.0/ee92df8e96c6907f3c851a987be3fd03d4b93b247e727b69a8e23ac94392a091. Subsequent calls will reuse this data.

0%| | 0/1 [00:00<?, ?it/s]

```
[ ]: from datasets import load_metric

metric = load_metric("accuracy")
```

Downloading builder script: 0%| | 0.00/1.41k [00:00<?, ?B/s]

```
[ ]: dataset
```

```
[ ]: DatasetDict({
    train: Dataset({
        features: ['image', 'label'],
        num_rows: 27000
    })
})
```

```
[ ]: example = dataset["train"][10]
example
```


```
[ ]: {'image': <PIL.JpegImagePlugin.JpegImageFile image mode=RGB size=64x64 at
0x7FF2F6277B10>,
     'label': 2}
```

```
[ ]: dataset["train"].features
```

```
[ ]: {'image': Image(decode=True, id=None),
     'label': ClassLabel(num_classes=10, names=['AnnualCrop', 'Forest',
'HerbaceousVegetation', 'Highway', 'Industrial', 'Pasture', 'PermanentCrop',
'Residential', 'River', 'SeaLake'], id=None)}
```


```
[ ]: example['image']
```

```
[ ]:
```



```
[ ]: example['image'].resize((200, 200))
```

```
[ ]:
```



```
[ ]: example['label']
```

```
[ ]: 2
```

```
[ ]: dataset["train"].features["label"]
```

```
[ ]: ClassLabel(num_classes=10, names=['AnnualCrop', 'Forest',  
    'HerbaceousVegetation', 'Highway', 'Industrial', 'Pasture', 'PermanentCrop',  
    'Residential', 'River', 'SeaLake'], id=None)
```

```
[ ]: labels = dataset["train"].features["label"].names  
    label2id, id2label = dict(), dict()  
    for i, label in enumerate(labels):  
        label2id[label] = i  
        id2label[i] = label  
  
    id2label[2]
```

```
[ ]: 'HerbaceousVegetation'
```

## 0.0.2 Preprocessing the data

```
[ ]: from transformers import AutoImageProcessor  
  
    image_processor = AutoImageProcessor.from_pretrained(model_checkpoint)  
    image_processor
```

```
Downloading: 0%|          | 0.00/255 [00:00<?, ?B/s]
```

```
[ ]: ViTFeatureExtractor {
    "do_normalize": true,
    "do_resize": true,
    "feature_extractor_type": "ViTFeatureExtractor",
    "image_mean": [
        0.485,
        0.456,
        0.406
    ],
    "image_std": [
        0.229,
        0.224,
        0.225
    ],
    "resample": 3,
    "size": 224
}
```

```
[ ]: from torchvision.transforms import (
    CenterCrop,
    Compose,
    Normalize,
    RandomHorizontalFlip,
    RandomResizedCrop,
    Resize,
    ToTensor,
)

normalize = Normalize(mean=image_processor.image_mean, std=image_processor.
    ↪image_std)
if "height" in image_processor.size:
    size = (image_processor.size["height"], image_processor.size["width"])
    crop_size = size
    max_size = None
elif "shortest_edge" in image_processor.size:
    size = image_processor.size["shortest_edge"]
    crop_size = (size, size)
    max_size = image_processor.size.get("longest_edge")

train_transforms = Compose(
    [
        RandomResizedCrop(crop_size),
        RandomHorizontalFlip(),
        ToTensor(),
        normalize,
    ]
)
```

```

val_transforms = Compose(
    [
        Resize(size),
        CenterCrop(crop_size),
        ToTensor(),
        normalize,
    ]
)

def preprocess_train(example_batch):
    """Apply train_transforms across a batch."""
    example_batch["pixel_values"] = [
        train_transforms(image.convert("RGB")) for image in
    ↪ example_batch["image"]
    ]
    return example_batch

def preprocess_val(example_batch):
    """Apply val_transforms across a batch."""
    example_batch["pixel_values"] = [val_transforms(image.convert("RGB")) for
    ↪ image in example_batch["image"]]
    return example_batch

```

```

[ ]: # split up training into training + validation
splits = dataset["train"].train_test_split(test_size=0.1)
train_ds = splits['train']
val_ds = splits['test']

```

```

[ ]: train_ds.set_transform(preprocess_train)
val_ds.set_transform(preprocess_val)

```

```

[ ]: train_ds[0]

```

```

[ ]: {'image': <PIL.JpegImagePlugin.JpegImageFile image mode=RGB size=64x64 at
0x7FF2EFFF0D90>,
      'label': 9,
      'pixel_values': tensor([[[-0.3541, -0.3541, -0.3541, ..., -0.3712, -0.3712,
-0.3712],
        [-0.3541, -0.3541, -0.3541, ..., -0.3712, -0.3712, -0.3712],
        [-0.3541, -0.3541, -0.3541, ..., -0.3712, -0.3712, -0.3712],
        ...,
        [-0.4397, -0.4397, -0.4397, ..., -0.4911, -0.4911, -0.4911],
        [-0.4397, -0.4397, -0.4397, ..., -0.4911, -0.4911, -0.4911],
        [-0.4397, -0.4397, -0.4397, ..., -0.4911, -0.4911, -0.4911]],
        [[-0.2500, -0.2500, -0.2500, ..., -0.2850, -0.2850, -0.2850],

```

```

        [-0.2500, -0.2500, -0.2500, ..., -0.2850, -0.2850, -0.2850],
        [-0.2500, -0.2500, -0.2500, ..., -0.2850, -0.2850, -0.2850],
        ...,
        [-0.3550, -0.3550, -0.3550, ..., -0.4076, -0.4076, -0.4076],
        [-0.3550, -0.3550, -0.3550, ..., -0.4076, -0.4076, -0.4076],
        [-0.3550, -0.3550, -0.3550, ..., -0.4076, -0.4076, -0.4076]],

[[ 0.1128, 0.1128, 0.1128, ..., 0.1651, 0.1651, 0.1651],
 [ 0.1128, 0.1128, 0.1128, ..., 0.1651, 0.1651, 0.1651],
 [ 0.1128, 0.1128, 0.1128, ..., 0.1651, 0.1651, 0.1651],
 ...,
 [ 0.0605, 0.0605, 0.0605, ..., 0.0082, 0.0082, 0.0082],
 [ 0.0605, 0.0605, 0.0605, ..., 0.0082, 0.0082, 0.0082],
 [ 0.0605, 0.0605, 0.0605, ..., 0.0082, 0.0082, 0.0082]]])}

```

### 0.0.3 Training the model

```

[ ]: from transformers import AutoModelForImageClassification, TrainingArguments, \
    Trainer

model = AutoModelForImageClassification.from_pretrained(
    model_checkpoint,
    label2id=label2id,
    id2label=id2label,
    ignore_mismatched_sizes = True, # provide this in case you're planning to
    fine-tune an already fine-tuned checkpoint
)

```

Downloading: 0%| | 0.00/70.1k [00:00<?, ?B/s]

Downloading: 0%| | 0.00/108M [00:00<?, ?B/s]

/usr/local/lib/python3.7/dist-packages/torch/functional.py:445: UserWarning:  
torch.meshgrid: in an upcoming release, it will be required to pass the indexing  
argument. (Triggered internally at

../aten/src/ATen/native/TensorShape.cpp:2157.)

```
    return _VF.meshgrid(tensors, **kwargs) # type: ignore[attr-defined]
```

Some weights of SwinForImageClassification were not initialized from the model  
checkpoint at microsoft/swin-tiny-patch4-window7-224 and are newly initialized  
because the shapes did not match:

- classifier.weight: found shape torch.Size([1000, 768]) in the checkpoint and  
torch.Size([10, 768]) in the model instantiated

- classifier.bias: found shape torch.Size([1000]) in the checkpoint and  
torch.Size([10]) in the model instantiated

You should probably TRAIN this model on a down-stream task to be able to use it  
for predictions and inference.

```
[ ]: model_name = model_checkpoint.split("/")[-1]
```

```
args = TrainingArguments(  
    f"{model_name}-finetuned-eurosat",  
    remove_unused_columns=False,  
    evaluation_strategy = "epoch",  
    save_strategy = "epoch",  
    learning_rate=5e-5,  
    per_device_train_batch_size=batch_size,  
    gradient_accumulation_steps=4,  
    per_device_eval_batch_size=batch_size,  
    num_train_epochs=3,  
    warmup_ratio=0.1,  
    logging_steps=10,  
    load_best_model_at_end=True,  
    metric_for_best_model="accuracy",  
    push_to_hub=True,  
)
```

```
[ ]: import numpy as np
```

```
# the compute_metrics function takes a Named Tuple as input:  
# predictions, which are the logits of the model as Numpy arrays,  
# and label_ids, which are the ground-truth labels as Numpy arrays.  
def compute_metrics(eval_pred):  
    """Computes accuracy on a batch of predictions"""  
    predictions = np.argmax(eval_pred.predictions, axis=1)  
    return metric.compute(predictions=predictions, references=eval_pred.  
↪label_ids)
```

```
[ ]: import torch
```

```
def collate_fn(examples):  
    pixel_values = torch.stack([example["pixel_values"] for example in  
↪examples])  
    labels = torch.tensor([example["label"] for example in examples])  
    return {"pixel_values": pixel_values, "labels": labels}
```

```
[ ]: trainer = Trainer(  
    model,  
    args,  
    train_dataset=train_ds,  
    eval_dataset=val_ds,  
    tokenizer=image_processor,  
    compute_metrics=compute_metrics,  
    data_collator=collate_fn,  
)
```



Cloning <https://huggingface.co/nielsr/swin-tiny-patch4-window7-224-finetuned-eurosat> into local empty directory.

```
[ ]: train_results = trainer.train()
      # rest is optional but nice to have
      trainer.save_model()
      trainer.log_metrics("train", train_results.metrics)
      trainer.save_metrics("train", train_results.metrics)
      trainer.save_state()
```

/usr/local/lib/python3.7/dist-packages/transformers/optimization.py:309:  
FutureWarning: This implementation of AdamW is deprecated and will be removed in  
a future version. Use the PyTorch implementation torch.optim.AdamW instead, or  
set `no\_deprecation\_warning=True` to disable this warning

```
FutureWarning,
***** Running training *****
    Num examples = 24300
    Num Epochs = 3
    Instantaneous batch size per device = 32
    Total train batch size (w. parallel, distributed & accumulation) = 128
    Gradient Accumulation steps = 4
    Total optimization steps = 570
```

<IPython.core.display.HTML object>

```
***** Running Evaluation *****
    Num examples = 2700
    Batch size = 32
Saving model checkpoint to swin-tiny-patch4-window7-224-finetuned-
eurosat/checkpoint-190
Configuration saved in swin-tiny-patch4-window7-224-finetuned-
eurosat/checkpoint-190/config.json
Model weights saved in swin-tiny-patch4-window7-224-finetuned-
eurosat/checkpoint-190/pytorch_model.bin
Feature extractor saved in swin-tiny-patch4-window7-224-finetuned-
eurosat/checkpoint-190/preprocessor_config.json
Feature extractor saved in swin-tiny-patch4-window7-224-finetuned-
eurosat/preprocessor_config.json
***** Running Evaluation *****
    Num examples = 2700
    Batch size = 32
Saving model checkpoint to swin-tiny-patch4-window7-224-finetuned-
eurosat/checkpoint-380
Configuration saved in swin-tiny-patch4-window7-224-finetuned-
eurosat/checkpoint-380/config.json
Model weights saved in swin-tiny-patch4-window7-224-finetuned-
eurosat/checkpoint-380/pytorch_model.bin
Feature extractor saved in swin-tiny-patch4-window7-224-finetuned-
eurosat/checkpoint-380/preprocessor_config.json
```

Feature extractor saved in swin-tiny-patch4-window7-224-finetuned-eurosat/preprocessor\_config.json

\*\*\*\*\* Running Evaluation \*\*\*\*\*

Num examples = 2700

Batch size = 32

Saving model checkpoint to swin-tiny-patch4-window7-224-finetuned-eurosat/checkpoint-570

Configuration saved in swin-tiny-patch4-window7-224-finetuned-eurosat/checkpoint-570/config.json

Model weights saved in swin-tiny-patch4-window7-224-finetuned-eurosat/checkpoint-570/pytorch\_model.bin

Feature extractor saved in swin-tiny-patch4-window7-224-finetuned-eurosat/checkpoint-570/preprocessor\_config.json

Feature extractor saved in swin-tiny-patch4-window7-224-finetuned-eurosat/preprocessor\_config.json

Training completed. Do not forget to share your model on [huggingface.co/models](https://huggingface.co/models)  
=)

Loading best model from swin-tiny-patch4-window7-224-finetuned-eurosat/checkpoint-570 (score: 0.9744444444444444).

Saving model checkpoint to swin-tiny-patch4-window7-224-finetuned-eurosat

Configuration saved in swin-tiny-patch4-window7-224-finetuned-eurosat/config.json

Model weights saved in swin-tiny-patch4-window7-224-finetuned-eurosat/pytorch\_model.bin

Feature extractor saved in swin-tiny-patch4-window7-224-finetuned-eurosat/preprocessor\_config.json

Saving model checkpoint to swin-tiny-patch4-window7-224-finetuned-eurosat

Configuration saved in swin-tiny-patch4-window7-224-finetuned-eurosat/config.json

Model weights saved in swin-tiny-patch4-window7-224-finetuned-eurosat/pytorch\_model.bin

Feature extractor saved in swin-tiny-patch4-window7-224-finetuned-eurosat/preprocessor\_config.json

Several commits (2) will be pushed upstream.

The progress bars may be unreliable.

Upload file pytorch\_model.bin: 0%| | 3.34k/105M [00:00<?, ?B/s]

Upload file runs/Apr12\_08-48-13\_9520b574893c/events.out.tfevents.1649753401.

9520b574893c.77.0: 24%|##4 ...

To <https://huggingface.co/nielsr/swin-tiny-patch4-window7-224-finetuned-eurosat-b46a767..6d6b8dc> main -> main

To <https://huggingface.co/nielsr/swin-tiny-patch4-window7-224-finetuned-eurosat>

```
6d6b8dc..25dd5d7  main -> main
```

```
***** train metrics *****
epoch                =          3.0
total_flos           = 1687935228GF
train_loss           =          0.3276
train_runtime        =    0:16:13.91
train_samples_per_second =          74.852
train_steps_per_second  =          0.585
```

We can check with the `evaluate` method that our `Trainer` did reload the best model properly (if it was not the last one):

```
[ ]: metrics = trainer.evaluate()
      # some nice to haves:
      trainer.log_metrics("eval", metrics)
      trainer.save_metrics("eval", metrics)
```

```
***** Running Evaluation *****
  Num examples = 2700
  Batch size = 32

<IPython.core.display.HTML object>
```

```
***** eval metrics *****
epoch                =          3.0
eval_accuracy        =          0.9744
eval_loss            =          0.0664
eval_runtime         =    0:00:16.12
eval_samples_per_second =          167.48
eval_steps_per_second  =          5.273
```

```
[ ]: trainer.push_to_hub()
```

```
Saving model checkpoint to swin-tiny-patch4-window7-224-finetuned-eurosat
Configuration saved in swin-tiny-patch4-window7-224-finetuned-
eurosat/config.json
```

```
Model weights saved in swin-tiny-patch4-window7-224-finetuned-
eurosat/pytorch_model.bin
```

```
Feature extractor saved in swin-tiny-patch4-window7-224-finetuned-
eurosat/preprocessor_config.json
```

```
Upload file runs/Apr12_08-48-13_9520b574893c/events.out.tfevents.1649754586.
9520b574893c.77.2: 100%|#####...
```

```
To https://huggingface.co/nielsr/swin-tiny-patch4-window7-224-finetuned-eurosat
25dd5d7..2164338  main -> main
```

```
[ ]: 'https://huggingface.co/nielsr/swin-tiny-patch4-window7-224-finetuned-eurosat/commit/2164338db59d40004286bc65800bfa50561ecd3d'
```

## 0.1 Inference

```
[ ]: from PIL import Image
import requests

url = 'https://huggingface.co/nielsr/convnext-tiny-finetuned-eurostat/resolve/
      ↪main/forest.png'
image = Image.open(requests.get(url, stream=True).raw)
image
```

```
[ ]:
```



```
[ ]: from transformers import AutoModelForImageClassification, AutoImageProcessor

repo_name = "nielsr/swin-tiny-patch4-window7-224-finetuned-eurosat"

image_processor = AutoImageProcessor.from_pretrained(repo_name)
model = AutoModelForImageClassification.from_pretrained(repo_name)
```

```
https://huggingface.co/nielsr/swin-tiny-patch4-window7-224-finetuned-eurosat/resolve/main/preprocessor_config.json not found in cache or
```

```

force_download set to True, downloading to
/root/.cache/huggingface/transformers/tmpqggthctf

Downloading: 0%|          | 0.00/240 [00:00<?, ?B/s]

storing https://huggingface.co/nielsr/swin-tiny-patch4-window7-224-finetuned-
eurosat/resolve/main/preprocessor_config.json in cache at /root/.cache/huggingfa
ce/transformers/7b742d61fc51f2ef5f81a75f80b26419c9f5bd86cc3022ed5784d09823f219f2
.e34548f8325ec440fcf4990d4a8dbbfd665397400e9a700766de032d2b45cf6b
creating metadata file for /root/.cache/huggingface/transformers/7b742d61fc51f2e
f5f81a75f80b26419c9f5bd86cc3022ed5784d09823f219f2.e34548f8325ec440fcf4990d4a8dbb
fd665397400e9a700766de032d2b45cf6b
loading feature extractor configuration file https://huggingface.co/nielsr/swin-
tiny-patch4-window7-224-finetuned-eurosat/resolve/main/preprocessor_config.json
from cache at /root/.cache/huggingface/transformers/7b742d61fc51f2ef5f81a75f80b2
6419c9f5bd86cc3022ed5784d09823f219f2.e34548f8325ec440fcf4990d4a8dbbfd665397400e9
a700766de032d2b45cf6b
Feature extractor ViTFeatureExtractor {
  "do_normalize": true,
  "do_resize": true,
  "feature_extractor_type": "ViTFeatureExtractor",
  "image_mean": [
    0.485,
    0.456,
    0.406
  ],
  "image_std": [
    0.229,
    0.224,
    0.225
  ],
  "resample": 3,
  "size": 224
}

```

```

https://huggingface.co/nielsr/swin-tiny-patch4-window7-224-finetuned-
eurosat/resolve/main/config.json not found in cache or force_download set to
True, downloading to /root/.cache/huggingface/transformers/tmpzdd89w3g

```

```

Downloading: 0%|          | 0.00/1.24k [00:00<?, ?B/s]

storing https://huggingface.co/nielsr/swin-tiny-patch4-window7-224-finetuned-
eurosat/resolve/main/config.json in cache at /root/.cache/huggingface/transforme
rs/83e4a1dea85e8e284e4da8ae1e3cf950c2c7e74d65a5a188049b3371fcd151bd.f1ed4852dd8f
4c3d0c565427607bc41fff51b58ac73a0970bec8456e5c64cea0
creating metadata file for /root/.cache/huggingface/transformers/83e4a1dea85e8e2
84e4da8ae1e3cf950c2c7e74d65a5a188049b3371fcd151bd.f1ed4852dd8f4c3d0c565427607bc4
1fff51b58ac73a0970bec8456e5c64cea0
loading configuration file https://huggingface.co/nielsr/swin-tiny-
patch4-window7-224-finetuned-eurosat/resolve/main/config.json from cache at /roo

```

t/.cache/huggingface/transformers/83e4a1dea85e8e284e4da8ae1e3cf950c2c7e74d65a5a188049b3371fcd151bd.f1ed4852dd8f4c3d0c565427607bc41fff51b58ac73a0970bec8456e5c64cea0

```
Model config SwinConfig {
  "_name_or_path": "nielsr/swin-tiny-patch4-window7-224-finetuned-eurosat",
  "architectures": [
    "SwinForImageClassification"
  ],
  "attention_probs_dropout_prob": 0.0,
  "depths": [
    2,
    2,
    6,
    2
  ],
  "drop_path_rate": 0.1,
  "embed_dim": 96,
  "encoder_stride": 32,
  "hidden_act": "gelu",
  "hidden_dropout_prob": 0.0,
  "hidden_size": 768,
  "id2label": {
    "0": "AnnualCrop",
    "1": "Forest",
    "2": "HerbaceousVegetation",
    "3": "Highway",
    "4": "Industrial",
    "5": "Pasture",
    "6": "PermanentCrop",
    "7": "Residential",
    "8": "River",
    "9": "SeaLake"
  },
  "image_size": 224,
  "initializer_range": 0.02,
  "label2id": {
    "AnnualCrop": 0,
    "Forest": 1,
    "HerbaceousVegetation": 2,
    "Highway": 3,
    "Industrial": 4,
    "Pasture": 5,
    "PermanentCrop": 6,
    "Residential": 7,
    "River": 8,
    "SeaLake": 9
  },
  "layer_norm_eps": 1e-05,
```

```

"mlp_ratio": 4.0,
"model_type": "swin",
"num_channels": 3,
"num_heads": [
    3,
    6,
    12,
    24
],
"num_layers": 4,
"patch_size": 4,
"path_norm": true,
"problem_type": "single_label_classification",
"qkv_bias": true,
"torch_dtype": "float32",
"transformers_version": "4.18.0",
"use_absolute_embeddings": false,
>window_size": 7
}

```

https://huggingface.co/nielsr/swin-tiny-patch4-window7-224-finetuned-eurosat/resolve/main/pytorch\_model.bin not found in cache or force\_download set to True, downloading to /root/.cache/huggingface/transformers/tmpkh0vdu53

Downloading: 0%| | 0.00/105M [00:00<?, ?B/s]

storing https://huggingface.co/nielsr/swin-tiny-patch4-window7-224-finetuned-eurosat/resolve/main/pytorch\_model.bin in cache at /root/.cache/huggingface/transformers/3daadbe0cabef18dc0e2232ae080d135a9d4ee6b1dc7675725ef38bedb990b81.818e63819e125637bd8a94f43b6899d1552f0b45884f1c28c458a5cb55dfa9e5

creating metadata file for /root/.cache/huggingface/transformers/3daadbe0cabef18dc0e2232ae080d135a9d4ee6b1dc7675725ef38bedb990b81.818e63819e125637bd8a94f43b6899d1552f0b45884f1c28c458a5cb55dfa9e5

loading weights file https://huggingface.co/nielsr/swin-tiny-patch4-window7-224-finetuned-eurosat/resolve/main/pytorch\_model.bin from cache at /root/.cache/huggingface/transformers/3daadbe0cabef18dc0e2232ae080d135a9d4ee6b1dc7675725ef38bedb990b81.818e63819e125637bd8a94f43b6899d1552f0b45884f1c28c458a5cb55dfa9e5

All model checkpoint weights were used when initializing SwinForImageClassification.

All the weights of SwinForImageClassification were initialized from the model checkpoint at nielsr/swin-tiny-patch4-window7-224-finetuned-eurosat. If your task is similar to the task the model of the checkpoint was trained on, you can already use SwinForImageClassification for predictions without further training.

```
[ ]: # prepare image for the model
      encoding = image_processor(image.convert("RGB"), return_tensors="pt")
      print(encoding.pixel_values.shape)
```

```
torch.Size([1, 3, 224, 224])
```

```
[ ]: import torch

      # forward pass
      with torch.no_grad():
          outputs = model(**encoding)
          logits = outputs.logits
```

```
[ ]: predicted_class_idx = logits.argmax(-1).item()
      print("Predicted class:", model.config.id2label[predicted_class_idx])
```

```
Predicted class: Forest
```